

Chapter 1

Proposed Design

Risk propagation is a message-passing algorithm that estimates an individual's infection risk by considering their demographics, symptoms, diagnosis, and contact with others. Formally, a *risk score* $s_t \in [0, 1]$ is a timestamped infection probability where $t \in \mathbb{N}$ is the time of its computation. Thus, an individual with a high risk score is likely to test positive for the infection and poses a significant health risk to others. There are two types of risk scores: *symptom scores*, or prior infection probabilities, which account for an individual's demographics, symptoms, and diagnosis [8, 33]; and *exposure scores*, or posterior infection probabilities, which incorporate the risk of direct and indirect contact with others.

Given their recent risk scores and contacts, an individual's exposure score is derived by marginalizing over the joint infection probability distribution. Naively computing this marginalization scales exponentially with the number of variables (i.e., individuals). To circumvent this intractability, the joint dis-

tribution is modeled as a factor graph, and an efficient message-passing procedure is employed to compute the marginal probabilities with a time complexity that scales linearly in the number of factor vertices (i.e., contacts).

Let $G = (X, F, E)$ be a *factor graph* where X is the set of variable vertices, F is the set of factor vertices, and E is the set of edges incident between them [27]. A *variable vertex* $x : \Omega \rightarrow \{0, 1\}$ is a random variable that represents the infection status of an individual, where the sample space is $\Omega = \{healthy, infected\}$ and

$$x(\omega) = \begin{cases} 0 & \text{if } \omega = healthy \\ 1 & \text{if } \omega = infected. \end{cases}$$

Thus, $p_i(x_i) = s_i$ is a risk score of the i -th individual. A *factor vertex* $f : X \times X \rightarrow [0, 1]$ defines the transmission of infection risk between two contacts. Specifically, contact between the i -th and j -th individual is represented by the factor vertex $f(x_i, x_j) = f_{ij}$, which is adjacent to the variable vertices x_i, x_j . Figure 1.1 depicts a factor graph that reflects the domain constraints.




Figure 1.1: A factor graph of 3 variable vertices and 2 factor vertices.

1.1 Synchronous Risk Propagation

Ayday, Yoo, and Halimi [3] first proposed risk propagation as a synchronous, iterative message-passing algorithm that uses the factor graph to compute exposure scores. The first input to RISK-PROPAGATION is the set family S , where

$$S_i = \{ s_t \mid \tau - t < T_s \} \in S \quad (1.1)$$

is the set of recent risk scores of the i -th individual. The second input to RISK-PROPAGATION is the contact set

$$C = \{ (i, j, t) \mid i \neq j, \tau - t < T_c \} \quad (1.2)$$

such that (i, j, t) is the *most recent* contact between the i -th and j -th individual that occurred from time t until at least time $t + \delta$, where $\delta \in \mathbb{N}$ is the *minimum contact duration*. Naturally, risk scores and contacts have finite relevance, so (1.1) and (1.2) are constrained by the *risk score expiry* $T_s \in \mathbb{N}$ and the *contact expiry* $T_c \in \mathbb{N}$, respectively. The *reference time* $\tau \in \mathbb{N}$ defines the relevance of the inputs and is assumed to be the time at which RISK-PROPAGATION is invoked. For notational simplicity in RISK-PROPAGATION, let X be a set. Then $\max X = 0$ if $X = \emptyset$.

1.1.1 Variable Messages

The current exposure score of the i -th individual is defined as $\max S_i$. Hence, a *variable message* $\mu_{ij}^{(n)}$ from the variable vertex x_i to the factor vertex f_{ij} during

the n -th iteration is the set of maximal risk scores $R_i^{(n-1)}$ from the previous $n - 1$ iterations that were not derived by f_{ij} . In this way, risk propagation is reminiscent of the max-sum algorithm; however, risk propagation aims to maximize *individual* marginal probabilities rather than the joint distribution [7, pp. 411–415].

1.1.2 Factor Messages

A *factor message* $\lambda_{ij}^{(n)}$ from the factor vertex f_{ij} to the variable vertex x_j during the n -th iteration is an exposure score of the j -th individual that is based on interacting with those at most $n - 1$ degrees separated from the i -th individual. This population is defined by the subgraph induced in G by

$$\{v \in X \cap F \setminus \{x_j, f_{ij}\} \mid d(x_i, v) \leq 2(n - 1)\},$$

where $d(u, v)$ is the shortest-path distance between the vertices u, v . The computation of a factor message assumes the following.

1. Contacts have a nondecreasing effect on an individual's exposure score.
2. A risk score s_t is *relevant* to the contact (i, j, t_{ij}) if $t < t_{ij} + \beta$, where $\beta \in \mathbb{N}$ is a *time buffer* that accounts for the incubation period, along with the delayed reporting of symptom scores and contacts. The time buffer is also known as the *(contact) tracing window* [39] and is similar to the *notification window* used by other contact tracing applications [29].
3. Risk transmission between contacts is incomplete. Thus, a risk score

decays exponentially along its transmission path in G at a rate of $\log \alpha$, where $\alpha \in (0, 1)$ is the *transmission rate*. Figure 1.2 visualizes this decay, assuming a transmission rate of $\alpha = 0.8$ [19].

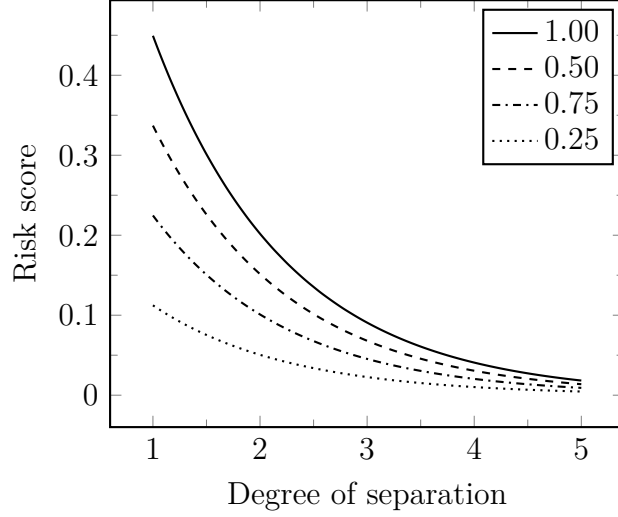


Figure 1.2: Exponential decay of risk scores.

To reiterate, a factor message $\lambda_{ij}^{(n)}$ is the maximum relevant risk score in the variable message $\mu_{ij}^{(n)}$ (or 0) that is scaled by the transmission rate α .

Ayday, Yoo, and Halimi [3] assume that the contact set C may contain (1) multiple contacts between the same two individuals and (2) invalid contacts, or those lasting less than δ time. However, these assumptions introduce unnecessary complexity. Regarding assumption 1, suppose the i -th and j -th individual come into contact m times such that $t_k < t_\ell$ for $1 \leq k < \ell \leq m$. Let Λ_k be the set of relevant risk scores, according to the contact time t_k , where

$$\Lambda_k = \{ \alpha s_t \mid s_t \in \mu_{ij}^{(n)}, t < t_k + \beta \}.$$

Then $\Lambda_k \subseteq \Lambda_\ell$ if and only if $\max \Lambda_k \leq \max \Lambda_\ell$. Therefore, only the most recent contact time t_m is required to compute the factor message $\lambda_{ij}^{(n)}$. With respect to assumption 2, there are two possibilities.

1. If an individual has at least one valid contact, then their exposure score is computed over the subgraph induced in G by their contacts that define the neighborhood N_i of the variable vertex x_i .
2. If an individual has no valid contacts, then their exposure score is $\max S_i$ or 0, if all of their previously computed risk scores have expired.

In either case, a set C containing only valid contacts implies fewer factor vertices and edges in the factor graph G . Consequently, the complexity of RISK-PROPAGATION is reduced by a constant factor since fewer messages must be computed.

1.1.3 Termination

To detect convergence, the normed difference between the current and previous exposure scores is compared to the threshold $\epsilon \in \mathbb{R}$. Note that $\mathbf{r}^{(n)}$ is the vector of exposure scores in the the n -th iteration such that $r_i^{(n)}$ is the i -th component of $\mathbf{r}^{(n)}$. The L_1 and L_∞ norms are sensible choices for detecting convergence. Ayday, Yoo, and Halimi [3] use the L_1 norm, which ensures that an individual's exposure score changed by at most ϵ after the penultimate iteration.

```

RISK-PROPAGATION( $S, C$ )
1:  $(X, F, E) \leftarrow \text{FACTOR-GRAPH}(C)$ 
2:  $n \leftarrow 1$ 
3: for each  $x_i \in X$ 
4:    $R_i^{(n-1)} \leftarrow \text{top } k \text{ of } S_i$ 
5:    $r_i^{(n-1)} \leftarrow \max R_i^{(n-1)}$ 
6:    $r_i^{(n)} \leftarrow \infty$ 
7: while  $\|\mathbf{r}^{(n)} - \mathbf{r}^{(n-1)}\| > \epsilon$ 
8:   for each  $\{x_i, f_{ij}\} \in E$ 
9:      $\mu_{ij}^{(n)} \leftarrow R_i^{(n-1)} \setminus \{\lambda_{ji}^{(\ell)} \mid \ell \in [1 \dots n-1]\}$ 
10:    for each  $\{x_i, f_{ij}\} \in E$ 
11:       $\lambda_{ij}^{(n)} \leftarrow \max \{\alpha s_t \mid s_t \in \mu_{ij}^{(n)}, t < t_{ij} + \beta\}$ 
12:    for each  $x_i \in X$ 
13:       $R_i^{(n)} \leftarrow \text{top } k \text{ of } \{\lambda_{ji}^{(n)} \mid f_{ij} \in N_i\}$ 
14:    for each  $x_i \in X$ 
15:       $r_i^{(n-1)} \leftarrow r_i^{(n)}$ 
16:       $r_i^{(n)} \leftarrow \max R_i^{(n)}$ 
17:     $n \leftarrow n + 1$ 
18: return  $\mathbf{r}^{(n)}$ 

```

1.2 Asynchronous Risk Propagation

While RISK-PROPAGATION offers proof of concept, it is not viable for real-world application. RISK-PROPAGATION is an *offline algorithm* [11], because it requires the contact and health information of all individuals as input. As

Ayday, Yoo, and Halimi [3] note, this centralization of personal data is not privacy-preserving. RISK-PROPAGATION also introduces communication overhead and computational redundancy since most exposure scores are unlikely to change across frequent invocations. To mitigate this inefficiency, Ayday et al. [4] suggest running RISK-PROPAGATION once or twice per day. Unfortunately, this cadence incurs substantial delay in updating individuals’ exposure scores. In the midst of a pandemic, timely information is essential for individual and collective health. To address the limitations of RISK-PROPAGATION, Ayday, Yoo, and Halimi [3] propose decentralizing the factor graph such that the processing entity associated with the i -th individual maintains the state of the i -th variable vertex and the neighboring factor vertices. Applying one-mode projection onto the variable vertices [45], Figure 1.3 illustrates how each entity corresponds to a portion of the factor graph. More generally, Ayday, Yoo, and Halimi [3] envision risk propagation as a decentralized communication protocol for informing individuals about their infection risk. Such a message-passing protocol naturally aligns with the *actor model*, a local model of concurrent computing that defines computation as patterns of message passing amongst computational entities called *actors* [2, 10, 15, 20, 21, 22]. Since communication is asynchronous in the actor model, risk propagation defined in this way is called *asynchronous risk propagation*.

1.2.1 Actor Behavior

An actor in the ShareTrace actor system corresponds to an individual. The CREATE-ACTOR operation [2] initializes an actor a with the following at-




Figure 1.3: One-mode projection onto the variable vertices in Figure 1.1.

tributes.

- *a.exposure*: the current exposure score of the individual. This attribute is either a symptom score, a risk score sent by another actor, or the null risk score (see NULL-RISK-SCORE).
- *a.contacts*: a *dictionary* (see ??) of contacts. In the context of an actor, a contact is a *proxy* [16] of the actor that represents an individual with which the individual represented by this actor was physically proximal. That is, if the i -th individual interacted with the j -th individual, then $a_i.contacts$ contains a contact c such that $c.key = c.name$ is a name of the j -th actor and $c.time$ is the most recent time of contact. This attribute extends the concept of *actor acquaintances* [2, 20, 21] to be time-varying.
- *a.scores*: a dictionary of exposure scores such that $s.key$ for an exposure score s is the time interval during which $a.exposure = s$. The null risk score is returned for queries in which the dictionary does not contain a risk score with a key that intersects the given query interval. Figure 1.4 depicts a hypothetical step function that $a.scores$ represents.

NULL-RISK-SCORE

```

1:  $s.value \leftarrow 0$ 
2:  $s.time \leftarrow 0$ 
3: return  $s$ 

```

CREATE-ACTOR

```

1:  $a.contacts \leftarrow \emptyset$ 
2:  $a.scores \leftarrow \emptyset$ 
3:  $a.exposure \leftarrow \text{NULL-RISK-SCORE}$ 
4: return  $a$ 

```




Figure 1.4: Hypothetical history of an individual's exposure score.

The interface of an actor is primarily defined by two types of messages: contacts and risk scores. Based on Section 1.1, let the *time to live* (TTL) of a message be the remaining time of its relevance. The reference time τ is assumed to be the current time.

RISK-SCORE-TTL(s)

1: **return** $T_s - (\tau - s.time)$

CONTACT-TTL(c)

1: **return** $T_c - (\tau - c.time)$

The HANDLE-RISK-SCORE operation defines how an actor behaves upon receiving a risk score. The key $s.key$ initially associated with the risk score s is the time interval during which it is relevant. For the dictionary $a.scores$, MERGE preserves the mapping invariant defined above such that risk scores are ordered first by value and then by time. Thus, $s.key \subseteq [s.time, s.time + T_s)$ for all exposure scores in $a.scores$. The UPDATE-EXPOSURE-SCORE operation is the online equivalent of updating an individual's exposure score in RISK-PROPAGATION.

HANDLE-RISK-SCORE(a, s)

```
1: if RISK-SCORE-TTL( $s$ ) > 0
2:    $s.key \leftarrow [s.time, s.time + T_s)$ 
3:   MERGE( $a.scores, s$ )
4:   UPDATE-EXPOSURE-SCORE( $a, s$ )
5:   for each  $c \in a.contacts$ 
6:     APPLY-RISK-SCORE( $a, c, s$ )
```

```

UPDATE-EXPOSURE-SCORE( $a, s$ )
1: if  $a.exposure.value < s.value$ 
2:    $a.exposure \leftarrow s$ 
3: else if  $RISK-SCORE-TTL(a.exposure) \leq 0$ 
4:    $a.exposure \leftarrow MAXIMUM(a.scores)$ 

```

For the moment, assume that APPLY-RISK-SCORE is the online equivalent of computing a factor message (see Section 1.1). Line 2 indicates that a copy s' of the risk score s is updated using the transmission rate α . The SEND operation enables actor communication [2].

```

APPLY-RISK-SCORE( $a, c, s$ )
1: if  $c.time + \beta > s.time$ 
2:    $s'.value \leftarrow \alpha \cdot s.value$ 
3:   SEND( $c.name, s'$ )

```

The problem with APPLY-RISK-SCORE is that it causes risk scores to propagate *ad infinitum*. Unlike RISK-PROPAGATION, a global convergence test is not available to terminate message passing, so it is necessary to define a local condition that determines if a risk score should be sent to another actor. Practically, “self-terminating” message-passing is necessary for asynchronous risk propagation to be scalable and cost-efficient.

The intent of sending a risk score to an actor is to update its exposure score. According to HANDLE-RISK-SCORE, it is only necessary to send an actor risk scores with values that exceed its current exposure score. Thus, an actor can associate a *send threshold* with a contact such that the target

actor only receives risk scores that exceed the threshold. To permit a trade-off between accuracy and efficiency, let the *send coefficient* $\gamma \in \mathbb{R}$ and *tolerance* $\epsilon \in \mathbb{R}$ be scaling factors that are applied to a risk score upon setting the send threshold.

SET-SEND-THRESHOLD(c, s)

- 1: $s'.value \leftarrow \gamma \cdot s.value + \epsilon$
- 2: $c.threshold \leftarrow s'$

The APPLY-RISK-SCORE operation that incorporates the concept of a send threshold is defined below. Assuming a finite number of actors, a positive send coefficient guarantees that a risk score has finite propagation.

APPLY-RISK-SCORE(a, c, s)

- 1: **if** $c.threshold.value < s.value$ **and** $c.time + \beta > s.time$
- 2: $s'.value \leftarrow \alpha \cdot s.value$
- 3: SET-SEND-THRESHOLD(c, s')
- 4: SEND($c.name, s'$)

The SET-SEND-THRESHOLD operation defines *how* the send threshold is updated, but not *when* it should be updated; the UPDATE-SEND-THRESHOLD operation encapsulates the latter. The second predicate on Line 1 stems from the fact that the send threshold is a risk score and thus subject to expiry. The first predicate, however, is more subtle and will be revisited shortly. The MAXIMUM-OLDER-THAN operation is the same as MAXIMUM, with the constraint that the key of the risk score intersects the interval $(-\infty, c.time + \beta)$. Thus, the returned risk score is always relevant to the contact. Consistent

with APPLY-RISK-SCORE, the risk score retrieved from $a.scores$ is scaled by the transmission rate and set as the new send threshold.

```

UPDATE-SEND-THRESHOLD( $a, c$ )
1: if  $c.threshold.value > 0$  and  $RISK-SCORE-TTL(c.threshold) \leq 0$ 
2:    $s \leftarrow MAXIMUM-OLDER-THAN(a.scores, c.time + \beta)$ 
3:    $s'.value \leftarrow \alpha \cdot s.value$ 
4:    $SET-SEND-THRESHOLD(c, s')$ 

```

The send threshold should be current when evaluating Line 1 of APPLY-RISK-SCORE above, so UPDATE-SEND-THRESHOLD is invoked beforehand.

```

APPLY-RISK-SCORE( $a, c, s$ )
1:  $UPDATE-SEND-THRESHOLD(a, c)$ 
2: if  $c.threshold.value < s.value$  and  $c.time + \beta > s.time$ 
3:    $s'.value \leftarrow \alpha \cdot s.value$ 
4:    $SET-SEND-THRESHOLD(c, s')$ 
5:    $SEND(c.name, s')$ 

```

Returning to the first predicate on Line 1 of UPDATE-SEND-THRESHOLD, the send threshold has a value of 0 when initially setting the send threshold as the null risk score and when no key in $a.scores$ intersects the query interval on (Line 2). Suppose the first predicate is omitted from Line 1. Given that UPDATE-SEND-THRESHOLD is the first statement in APPLY-RISK-SCORE, it is possible that the send threshold is set prior to sending the target actor a relevant risk score. In the worst case, this prevents *all* risk scores from being sent to that actor. As a result, that individual would not receive risk scores

that may impact their exposure score. Therefore, for correct message-passing behavior, the send threshold is updated only when its value is nonzero.

The aforementioned refinements to APPLY-RISK-SCORE have focused on ensuring that message passing terminates and behaves correctly over time. To conclude the definition of APPLY-RISK-SCORE, a message-passing optimization, called *sender-side aggregation* [32], is incorporated. Over a given period of time, an actor may receive several risk scores that are subsequently sent to multiple target actors. Rather than sending multiple risk scores, it would be more efficient to just send the final risk score. As a heuristic, APPLY-RISK-SCORE can be modified so that a contact “buffers” a risk score that the target actor should receive.

```

APPLY-RISK-SCORE( $a, c, s$ )
1: UPDATE-SEND-THRESHOLD( $a, c$ )
2: if  $c.threshold.value < s.value$  and  $c.time + \beta > s.time$ 
3:    $s'.value \leftarrow \alpha \cdot s.value$ 
4:   SET-SEND-THRESHOLD( $c, s'$ )
5:    $c.buffered \leftarrow s'$ 

```

When the actor receives a periodic *flush timeout* message, all contacts are “flushed” by sending the buffered risk scores to the target actors. For convenience, the HANDLE-FLUSH-TIMEOUT operation also removes all expired contacts from $a.contacts$.

HANDLE-FLUSH-TIMEOUT(a)

```
1: for each  $c \in a.contacts$ 
2:   if  $c.buffered \neq \text{NIL}$ 
3:     SEND( $c.name, c.buffered$ )
4:      $c.buffered \leftarrow \text{NIL}$ 
5:   if CONTACT-TTL( $c$ )  $\leq 0$ 
6:     DELETE( $a.contacts, c$ )
```

The HANDLE-CONTACT operation concludes this section on actor behavior. Similar to HANDLE-RISK-SCORE, expired contacts are not processed. The MERGE operation for $a.contacts$ differs from its usage with $a.scores$. That is, if a contact with the same key already exists, its contact time is updated to that of the newer contact; all other state of the previous contact is maintained. A risk score from $a.scores$ is also applied to the contact to ensure that, if the actor receives no other risk score before the contact expires, the target actor is sent at least one relevant risk score.

HANDLE-CONTACT(a, c)

```
1: if CONTACT-TTL( $c$ )  $> 0$ 
2:    $c.threshold \leftarrow \text{NULL-RISK-SCORE}$ 
3:    $c.buffered \leftarrow \text{NIL}$ 
4:    $c.key \leftarrow c.name$ 
5:   MERGE( $a.contacts, c$ )
6:    $s \leftarrow \text{MAXIMUM-OLDER-THAN}(a.scores, c.time + \beta)$ 
7:   APPLY-RISK-SCORE( $a, c, s$ )
```


1.2.2 Message Reachability

The ShareTrace actor system is a *contact network*, a type of *temporal network* in which vertices represent individuals and edges indicate contact between them [23, 24]. The contact set defined in (1.2) is the *contact sequence* representation of a contact network [24]. Contact networks are typically used in epidemiological studies that aim to model and analyze the spreading dynamics of infection [12, 14, 26, 30, 37, 40, 46]. The ShareTrace actor network, however, models the spreading of infection *risk*. Holme and Saramäki [24] note that the transmission graph proposed by Riolo, Koopman, and Chick [40] “cannot handle edges where one node manages to not catch the disease.” By framing the spreading phenomenon as continuous, rather than discrete, it is possible to model partial transmission.

A primitive of temporal reachability analysis is the *time-respecting path*: a contiguous sequence of contacts with nondecreasing time. Thus, vertex v is *temporally reachable* from vertex u if there exists a time-respecting path from u to v [34]. The following derivatives of a time-respecting path help quantify reachability in temporal networks [24].

- *influence set* $I(v)$: vertices that v can reach by a time-respecting path.
- *source set* $S(v)$: vertices that can reach v by a time-respecting path.
- *reachability ratio* $\rho(G)$: the average influence set cardinality in G .

Generally, a message-passing algorithm specifies constraints that determine when and what messages are sent between vertices. Even if operating on a temporal network, those constraints may be more or less strict than requiring

temporal reachability. As a dynamic process, message passing on a time-varying network requires a broader definition of reachability that accounts for network topology *and* message-passing semantics [5].

Formally, the *reachability* of a message m from vertex u to vertex v is the number of edges along the shortest path P that satisfy the message-passing constraints,

$$r(u, v) = \sum_{(i,j) \in P} f(u, i, j, v),$$

where

$$f(u, i, j, v) = \begin{cases} 1 & \text{if all constraints are satisfied} \\ 0 & \text{otherwise.} \end{cases}$$

Vertex v is *reachable* from vertex u for the message m if there exists a shortest path P such that $r(u, v) = |P|$. The *reachability* of a message m from vertex u is

$$r(u) = \max_{v \in V} r(u, v). \quad (1.3)$$

Measures of temporal reachability can be extended to message reachability,

$$I(u) = \{ v \in V \mid r(u, v) = |P| \}$$

$$S(v) = \{ u \in V \mid r(u, v) = |P| \}$$

$$\rho(G, M) = \sum_{v \in V, m \in M} |I(v)| \cdot |V|^{-1},$$

where M is the set of messages associated with the vertices V .

Asynchronous Risk Propagation

Let P be the set of contact edges along the shortest path from actor u to actor v such that the actors are enumerated $1, \dots, |P|$. Let (s_u, t_u) be a symptom score of actor u . Let s_{ij} be the value of the risk score that is associated with the send threshold of the i -th actor for the j -th actor. Finally, let t_{ij} be the most recent contact time between the i -th and j -th individual. Then message reachability for asynchronous risk propagation is defined as

$$r(u, v) = \sum_{(i,j) \in P} [\alpha^i s_u > \gamma \alpha s_{ij} + \epsilon] \cdot [t_u < t_{ij} + \beta], \quad (1.4)$$

where $[\cdot]$ is the Iverson bracket¹.

By relaxing the temporal constraint in (1.4), an upper bound on the reachability of a symptom score can be defined. Reversing the inequality of the first term in (1.4) and solving for i ,

$$\hat{r}(u, v) \leq \begin{cases} 0 & \text{if } s_u = 0 \\ |P| & \text{if } s_v = 0 \\ \log_{\alpha} \frac{\gamma \alpha s_v + \epsilon}{s_u} & \text{otherwise,} \end{cases} \quad (1.5)$$

where $\gamma \alpha s_v + \epsilon$ is the send threshold of the actor that precedes actor v along the path P . Assuming a transmission rate of $\alpha = 0.8$, a send coefficient of $\gamma = 1$, and a tolerance of $\epsilon = 0$, Figure 1.5 visualizes (1.5).

¹The *Iverson bracket* $[x]$ is the indicator function of the set values for which x is true.




Figure 1.5: Message reachability for asynchronous risk propagation. Contour lines are shown for integral reachability values, indicating the maximum send threshold that permits sending the risk score.

1.3 System Design

With the algorithmic foundation of ShareTrace established, it is now possible to discuss the aspects of system design. To a varying extent, prior work on ShareTrace [3, 4, ??] has explored the practical considerations for deployment. This work, however, provides additional detail by contextualizing ShareTrace as an application of *mobile crowdsensing* (MCS), a “sensing paradigm that empowers ordinary citizens to contribute data sensed or generated from their mobile devices” that is aggregated “for crowd intelligence extraction and human-centric service delivery” [18]. To offer a clear comparison with other MCS applications, this section uses the classification criteria proposed by Capponi et al. [9]. Note, this section assumes the usage of asynchronous risk propagation that is described in Section 1.2.

1.3.1 Application Layer

Task Scheduling

Proactive scheduling allows users to decide when and where they contribute sensing data, while *reactive scheduling* requires that “a user receives a request, and upon acceptance, accomplishes a task” [9]. ShareTrace follows proactive scheduling where the sensing task is to detect proximal interactions with other individuals. Naturally, the scheduling of this task is at the discretion of ShareTrace users.

Task Assignment

Centralized assignment assumes that “a central authority distributes tasks among users.” Conversely, with *decentralized assignment*, “each participant becomes an authority and can either perform the task or forward it to other users” [9]. The latter aligns with ShareTrace since each user is responsible for their own interactions.

Task Execution

With *single-task execution*, MCS applications assign one type of task to users, while *multi-task execution* assigns multiple types of tasks [9]. ShareTrace involves the single task of sensing proximal interactions to derive an individual’s infection risk.

User Recruitment

Volunteer-based recruitment is when citizens can “join a sensing campaign for personal interests...or willingness to help the community,” while *incentive-based recruitment* promotes participation and offers control over the recruitment rate...These strategies are not mutually exclusive and users can first volunteer and then be rewarded for quality execution of sensing tasks” [9]. ShareTrace follows volunteer-based recruitment, though incentive mechanisms would be an important consideration for widespread adoption [1, 36].

User Selection

User-centric selection is when “contributions depend only on participants['] willingness to sense and report data to the central collector, which is not responsible for choosing them.” *Platform-centric selection* is “when the central authority directly decides data contributors...Platform centric decisions are taken according to the utility of data to accomplish the targets of the campaign” [9]. ShareTrace employs user-centric selection, because the purpose of the application is to passively sense in-person interactions and provide individuals with knowledge of their infection risk. However, ShareTrace does not require the presence of a central collector, such as a government health authority.

User Type

A *contributor* “reports data to the MCS platform with no interest in the results of the sensing campaign” and “are typically driven by incentives or by

the desire to help the scientific or civil communities.” A *consumer* joins “a service to obtain information about certain application scenario[s] and have a direct interest in the results of the sensing campaign” [9]. ShareTrace users could be contributors or consumers but would likelier be the latter.

1.3.2 Data Layer

Data Storage

Centralized storage involves data being “stored and maintained in a single location, which is usually a database made available in the cloud. This approach is typically employed when significant processing or data aggregation is required.” *Distributed storage* “is typically employed for delay-tolerant applications, i.e., when users are allowed to deliver data with a delayed reporting” [9].

ShareTrace relies on distributed—more specifically, decentralized—storage. Prior work [3, 4, ??] suggests using Dataswyft Personal Data Accounts², which provide a data-oriented interface to self-sovereign identity [38, 41]. However, personal data stores alone cannot provide individuals true controllership over their personal data and have historically not demonstrated successful adoption in the consumer market [35]. Regardless of whether message-passing between actors is decentralized, so long as actor computation is decentralized on mobile devices, using local storage is the simplistic approach.

To allow asynchronous message passing between actors in the ShareTrace actor system (i.e., to be delay-tolerant), the underlying messaging system must

²<https://dataswyft.io>

persist the messages sent between actors, at least until they expire. While further research on alternative decentralized technologies is needed [25, 44], the InterPlanetary FileSystem (IPFS)³ [6, 42, 43], and the various higher-level storage abstractions (e.g., OrbitDB⁴ and web3.storage⁵) that have been developed with it, offers a promising solution. Additionally, investigating the feasibility of implementing secure publish-subscribe systems [13] with IPFS could provide robust security and privacy guarantees.

Data Format

Structured data is standardized and readily analyzable, while *unstructured data* requires significant processing before it can be used [9]. ShareTrace data, such as reported symptoms, biometric data, contact identifiers, and risk scores, is structured.

Data Dimensionality

Single-dimension data typically occurs when a single sensor is used, while *multi-dimensional data* arises with the use of multiple sensors. Prior work on ShareTrace assumes single-dimensional data: contact identifiers. Though, integrating biometric sensor data into the calculation of symptom scores would classify ShareTrace as using multi-dimensional data.

³<https://ipfs.tech>

⁴<https://orbitdb.org>

⁵<https://web3.storage>

Data Pre-processing

Raw data output implies that no modification is made to the sensed data. *Filtering and denoising* entail “removing irrelevant and redundant data. In addition, they help to aggregate and make sense of data while reducing at the same time the volume to be stored” [9]. ShareTrace should aggregate contact identifiers over time in order to determine which encounters were sustained (e.g., lasting at least 15 minutes). Moreover, if biometric sensors are also incorporated into the symptom score calculation, then signal processing would likely be required.

Data Analytics

Machine learning (ML) and data mining analytics “aim to infer information, identify patterns, or predict future trends” and are not real-time. On the contrary, *real-time analytics* consist of “examining collected data as soon as it is produced by the contributors” [9]. While the message passing that occurs during risk propagation may take place in near real-time, this is not a requirement. Thus, ShareTrace more closely aligns with the former category since it aims to predict an individual’s infection risk.

Data Post-processing

Statistical post-processing “aims at inferring proportions given quantitative examples of the input data.” *Prediction post-processing* aims to determine “future outcomes from a set of possibilities when given new input in the system” [9]. ShareTrace applies predictive post-processing via risk propagation.

1.3.3 Communication Layer

Infrastructured Technology

Cellular connectivity “is typically required from sensing campaign[s] that perform real-time monitoring and require data reception as soon as it is sensed.”

WLAN “is used mainly when sensing organizers do not specify any preferred reporting technologies or when the application domain permits to send data” at “a certain amount of time after the sensing process” [9]. Infrastructured technology is also referred to as the *infrastructured transmission paradigm* [31]. ShareTrace is delay-tolerant; thus, it can relay on WLAN infrastructure.

Infrastructureless Technology

Infrastructureless technologies “consists of device-to-device (D2D) communications that do not require any infrastructure...but rather allow devices in the vicinity to communicate directly.” Technologies include *WiFi-Direct*, *LTE-Direct*, and *Bluetooth* [9]. Infrastructureless technology is also called the *opportunistic transmission paradigm* [31]. Because of its sufficient short-range accuracy and energy efficiency, ShareTrace uses Bluetooth Low Energy (BLE) to transmit and receive ephemeral identifiers.

Upload mode

With *relay uploading*, “data is delivered as soon as collected.” *Store-and-forward* “is typically used in delay-tolerant applications when campaigns do not need to receive data in real-time” [9]. ShareTrace is delay-tolerant, so it

uses store-and-forward uploading.

Reporting Methodology

Individualized sensing is “when each user accomplishes the requested task individually and without interaction with other participants.” *Collaborative sensing* is when “users communicate with each other, exchange data[,] and help themselves in accomplishing a task...Users are typically grouped and exchange data exploiting short-range communication technologies” [9]. The methodology of sensing is similar to the *sensing scale* which is typically dichotomized as *personal* [17, 28] (i.e., individualized) and *community* [17] or *group* [28] (i.e., collaborative). ShareTrace includes elements of both individualized and collaborative sensing. With respect to the former, each user may individually use biometric sensors to calculate their symptom score. At the same time, ShareTrace users collaboratively sense each other in order to estimate their exposure score. The latter aspect of ShareTrace is a requirement, so overall ShareTrace primarily applies collaborative sensing.

Reporting Timing

Timing is based on whether devices “should sense in the same period or not.” *Synchronous timing* “includes cases in which users start and accomplish at the same time the sensing task. For synchronization purposes, participants communicate with each other.” *Asynchronous timing* occurs “when users perform sensing activity not in time synchronization with other users” [9]. ShareTrace requires synchronous timing, because contact sensing inherently requires syn-

chronous communication between mobile devices.

1.3.4 Sensing Layer

Sensor Deployment

Dedicated deployment involves using “non-embedded sensing elements,” typically for a specific task. *Non-dedicated deployment* utilizes sensors that “do not require to be paired with other devices for data delivery but exploit the communication capabilities of mobile devices” [9]. ShareTrace relies on non-dedicated deployment since it relies on BLE sensors that are ubiquitous in modern-day mobile devices.

Sensor Activity

Always-on sensors “are required to accomplish mobile devices['] basic functionalities, such as detection of rotation and acceleration...Activity recognition [i.e., context awareness]...is a very important feature that accelerometers enable.” *On-demand sensors* “need to be switched on by users or exploiting an application running in the background. Typically, they serve more complex applications than always-on sensors and consume a higher amount of energy” [9]. ShareTrace uses on-demand sensors (i.e., BLE), which run in the background and is not required for basic device functionality. ShareTrace may utilize always-on sensors, such as the accelerometer, to contextually activate the on-demand sensor and improve energy efficiency.

Data Acquisition

Homogeneous acquisition “involves only one type of data and it does not change from one user to another one,” while *heterogeneous acquisition* “involves different data types usually sampled from several sensors” [9]. In the event that ShareTrace only senses contact identifiers, it would be considered to use homogenous acquisition. If however, optional biometric sensors could be integrated into symptom score calculation, then ShareTrace would follow heterogenous acquisition.

Sampling Frequency

Continuous sensing “indicates tasks that are accomplished regularly and independently [of] the context of the smartphone or the user[’s] activities.” *Event-based sensing* is “data collection [that] starts after a certain event has occurred. In this context, an event can be seen as an active action from a user or the central collector, but also a given context awareness” [9]. ShareTrace continuous senses contact identifiers. However, in an effort to conserve energy, motion sensors could be used to selectively activate sensing when the user is carrying their mobile device.

Sensing Responsibility

When the *mobile device* is responsible, “devices or users take sampling decisions locally and independently from the central authority...It is often necessary to detect the context [of the] smartphones and wearable devices.” When the *central collector* is responsible, they make “decisions about sensing and

communicate them to the mobile devices” [9]. Given the human-centric and decentralized nature of the ShareTrace sensing task, mobile devices are responsible.

User involvement

Participatory involvement “requires active actions from users, who are explicitly asked to perform specific tasks. They are responsible to consciously meet the application requests by deciding when, what, where, and how to perform sensing tasks.” With *Opportunistic involvement*, “users...only declare their interest in joining a campaign and providing their sensors as a service...The smartphone itself is context-aware and makes decisions to sense and store data, automatically determining when its context matches the requirements of an application” [9]. Earlier works on MCS refer to user involvement as the *sensing paradigm* [17, 28, 31]. ShareTrace applies opportunistic involvement since users do not need to take explicit action beyond installing the application on their mobile device. As stated previously, ShareTrace may also integrate on-demand sensors in order to contextually decide when to sense nearby users.

Application	Task	Scheduling	Proactive Reactive	●
		Assignment	Centralized Decentralized	●
		Execution	Single task Multi-tasking	●
	User	Recruitment	Voluntary Incentivized	● ○
		Selection	Platform-centric User-centric	●
		Type	Consumer Contributor	● ○
Data	Management	Storage	Centralized Distributed	●
		Format	Structured Unstructured	●
		Dimension	Single dimension Multi-dimensional	● ○
	Processing	Pre-processing	Raw data Filtering and denoising	●
		Analytics	ML and data mining Real-time	●
		Post-processing	Statistical Prediction	●
Communication	Technologies	Infrastructured	Cellular WLAN	●
		Infrastructureless	LTE-Direct WiFi-Direct Bluetooth	●
	Reporting	Upload mode	Relay Store and forward	●
		Methodology	Individual Collaborative	○ ●
		Timing	Synchronous Asynchronous	●
Sensing	Elements	Deployment	Dedicated Non-dedicated	●
		Activity	Always-on On-demand	● ○
		Acquisition	Homogeneous Heterogeneous	● ○
	Sampling	Frequency	Continuous Event-based	● ○
		Responsibility	Mobile device Central collector	●
		User involvement	Participatory Opportunistic	●

Table 1.1: Classification of ShareTrace as a mobile crowdsensing (MCS) application. This table follows the four-layered architecture of MCS applications proposed by Capponi et al. [9]. Proposed design (●); plausible alternative (○).

Bibliography

- [1] Saleh Afroogh et al. “Tracing app technology: An ethical review in the COVID-19 era and directions for post-COVID-19”. In: *Ethics and Information Technology* 24.3 (2022). DOI: 10.1007/s10676-022-09659-6 (cit. on p. 22).
- [2] Gul Agha. “ACTORS: A model of concurrent computation in distributed systems”. PhD thesis. MIT, 1985. URL: <https://hdl.handle.net/1721.1/6952> (cit. on pp. 8, 9, 12).
- [3] Erman Ayday, Youngjin Yoo, and Anisa Halimi. “ShareTrace: An iterative message passing algorithm for efficient and effective disease risk assessment on an interaction graph”. In: *Proceedings of the 12th ACM Conference on Bioinformatics, Computational Biology, and Health Informatics*. 2021. DOI: 10.1145/3459930.3469553 (cit. on pp. 3, 5, 6, 8, 20, 23).
- [4] Erman Ayday et al. *ShareTrace: A smart privacy-preserving contact tracing solution by architectural design during an epidemic*. White paper. Case Western Reserve University, 2020 (cit. on pp. 8, 20, 23).

- [5] Alain Barrat and Ciro Cattuto. “Temporal networks of face-to-face human interactions”. In: *Temporal Networks*. Springer, 2013. DOI: 10.1007/978-3-642-36461-7_10 (cit. on p. 18).
- [6] Juan Benet. *IPFS - content addressed, versioned, P2P file system*. 2014. arXiv: 1407.3561 [cs.NI] (cit. on p. 24).
- [7] Christopher M. Bishop. “Pattern recognition and machine learning”. In: *Information Science and Statistics*. Springer, 2006 (cit. on p. 4).
- [8] Mark Briers, Marcos Charalambides, and Chris Holmes. *Risk scoring calculation for the current NHSx contact tracing app*. 2020. arXiv: 2005.11057 [cs.CY] (cit. on p. 1).
- [9] Andrea Capponi et al. “A survey on mobile crowdsensing systems: Challenges, solutions, and opportunities”. In: *IEEE Communication Surveys & Tutorials* 21.3 (2019), pp. 2419–2465. DOI: 10.1109/comst.2019.2914030 (cit. on pp. 20–31).
- [10] William Clinger. “Foundations of actor semantics”. PhD thesis. MIT, 1981. URL: <https://hdl.handle.net/1721.1/6935> (cit. on p. 8).
- [11] Thomas H. Cormen et al. *Introduction to algorithms*. The MIT Press, 2022 (cit. on p. 7).
- [12] Meggan E. Craft. “Infectious disease transmission and contact networks in wildlife and livestock”. In: *Philosophical Transactions of the Royal Society B* 370.1669 (2015). DOI: 10.1098/rstb.2014.0107 (cit. on p. 17).

- [13] Shujie Cui et al. “Collusion defender: Preserving subscribers’ privacy in publish and subscribe systems”. In: *IEEE Transactions on Dependable and Secure Computing* 18.3 (2021), pp. 1051–1064. DOI: 10.1109/tdsc.2019.2898827 (cit. on p. 24).
- [14] Leon Danon et al. “Networks and the epidemiology of infectious disease”. In: *Interdisciplinary Perspectives on Infectious Diseases* 2011 (2011). DOI: 10.1155/2011/284909 (cit. on p. 17).
- [15] Joeri De Koster, Tom Van Cutsem, and Wolfgang De Meuter. “43 years of actors: A taxonomy of actor models and their key properties”. In: *Proceedings of the 6th International Workshop on Programming Based on Actors, Agents, and Decentralized Control*. 2016, pp. 31–40. DOI: 10.1145/3001886.3001890 (cit. on p. 8).
- [16] Erich Gamma et al. *Design patterns: Elements of reusable object-oriented software*. Addison-Wesley, 1995 (cit. on p. 9).
- [17] Raghu K. Ganti, Fan Ye, and Hui Lei. “Mobile crowdsensing: Current state and future challenges”. In: *IEEE Communications Magazine* 49.11 (2011), pp. 32–39. DOI: 10.1109/mcom.2011.6069707 (cit. on pp. 27, 30).
- [18] Bin Guo et al. “Mobile crowd sensing and computing: The review of an emerging human-powered sensing paradigm”. In: *ACM Computing Surveys* 48.1 (2015), pp. 1–31. DOI: 10.1145/2794400 (cit. on p. 20).
- [19] Lea Hamner et al. “High SARS-CoV-2 attack rate following exposure at a choir practice – Skagit County, Washington, March 2020”. In: *Mor-*

- bidity and Mortality Weekly Report* 69.19 (2020). DOI: 10.15585/mmwr.mm6919e6 (cit. on p. 5).
- [20] Carl Hewitt. “Viewing control structures as patterns of passing messages”. In: *Artificial Intelligence* 8.3 (1977), pp. 323–364. DOI: 10.1016/0004-3702(77)90033-9 (cit. on pp. 8, 9).
 - [21] Carl Hewitt and Henry Baker. *Laws for communicating parallel processes*. Working paper. MIT Artificial Intelligence Laboratory, 1977 (cit. on pp. 8, 9).
 - [22] Carl Hewitt, Peter Bishop, and Richard Steiger. “A universal modular ACTOR formalism for artificial intelligence”. In: *Proceedings of the 3rd International Joint Conference on Artificial Intelligence*. 1973, pp. 235–245. URL: <https://dl.acm.org/doi/10.5555/1624775.1624804> (cit. on p. 8).
 - [23] Petter Holme. “Modern temporal network theory: A colloquium”. In: *The European Physical Journal B* 88.9 (2015). DOI: 10.1140/epjb/e2015-60657-4 (cit. on p. 17).
 - [24] Petter Holme and Jari Saramäki. “Temporal networks”. In: *Physics Reports* 519.3 (2012). DOI: 10.1016/j.physrep.2012.03.001 (cit. on p. 17).
 - [25] Navin Keizer et al. “A survey on content retrieval on the decentralised web”. In: *ACM Computing Surveys* 56.8 (2024). DOI: 10.1145/3649132 (cit. on p. 24).

- [26] Andreas Koher et al. “Contact-based model for epidemic spreading on temporal networks”. In: *Physical Review X* 9.3 (2019). DOI: 10.1103/physrevx.9.031017 (cit. on p. 17).
- [27] Frank R. Kschischang, Brendan J. Frey, and Hans A. Loeliger. “Factor graphs and the sum-product algorithm”. In: *IEEE Transactions on Information Theory* 47.2 (2001). DOI: 10.1109/18.910572 (cit. on p. 2).
- [28] Nicholas D. Lane et al. “A survey of mobile phone sensing”. In: *IEEE Communications Magazine* 48.9 (2010), pp. 140–150. DOI: 10.1109/mcom.2010.5560598 (cit. on pp. 27, 30).
- [29] Trystan Leng et al. “The effect of notification window length on the epidemiological impact of COVID-19 contact tracing mobile applications”. In: *Communications Medicine* 2.1 (2022). DOI: 10.1038/s43856-022-00143-2 (cit. on p. 4).
- [30] Andrey Y. Lokhov et al. “Inferring the origin of an epidemic with a dynamic message-passing algorithm”. In: *Physical Review E* 90.1 (2014). DOI: 10.1103/physreve.90.012801 (cit. on p. 17).
- [31] Huadong Ma, Dong Zhao, and Peiyan Yuan. “Opportunities in mobile crowd sensing”. In: *IEEE Communications Magazine* 52.8 (2014), pp. 29–35. DOI: 10.1109/mcom.2014.6871666 (cit. on pp. 26, 30).
- [32] Robert McCune, Tim Weninger, and Greg Madey. “Thinking like a vertex: A survey of vertex-centric frameworks for large-scale distributed graph processing”. In: *ACM Computing Surveys* 48.2 (2015). DOI: 10.1145/2818185 (cit. on p. 15).

- [33] Cristina Menni et al. “Real-time tracking of self-reported symptoms to predict potential COVID-19”. In: *Nature Medicine* 26.7 (2020). DOI: 10.1038/s41591-020-0916-2 (cit. on p. 1).
- [34] James Moody. “The importance of relationship timing for diffusion”. In: *Social Forces* 81.1 (2002). URL: <https://www.jstor.org/stable/3086526> (cit. on p. 17).
- [35] Arvind Narayanan et al. *A critical look at decentralized personal data architectures*. 2012. arXiv: 1202.4503 [cs.CY] (cit. on p. 23).
- [36] Kiemute Oyibo et al. “Factors influencing the adoption of contact tracing applications: Systematic review and recommendations”. In: *Frontiers in Digital Health* 4 (2022). DOI: 10.3389/fdgth.2022.862466 (cit. on p. 22).
- [37] Romualdo Pastor-Satorras et al. “Epidemic processes in complex networks”. In: *Reviews of Modern Physics* 87.3 (2015). DOI: 10.1103/revmodphys.87.925 (cit. on p. 17).
- [38] Alex Preukschat and Drummond Reed. *Self-sovereign identity: Decentralized digital identity and verifiable credentials*. Manning, 2021 (cit. on p. 23).
- [39] Joren Raymenants et al. “Empirical evidence on the efficiency of backward contact tracing in COVID-19”. In: *Nature Communications* 13.1 (2022). DOI: 10.1038/s41467-022-32531-6 (cit. on p. 4).
- [40] Christopher S. Riolo, James S. Koopman, and Stephen E. Chick. “Methods and measures for the description of epidemiologic contact networks”.

- In: *Journal of Urban Health* 78.3 (2001). DOI: 10.1093/jurban/78.3.446 (cit. on p. 17).
- [41] Frederico Schardong and Ricardo Custódio. “Self-sovereign identity: A systematic review, mapping and taxonomy”. In: *Sensors* 22.15 (2022). DOI: 10.3390/s22155641 (cit. on p. 23).
 - [42] Ruizhe Shi et al. “A closer look into IPFS: Accessibility, content, and performance”. In: *Proceedings of the ACM on Measurement and Analysis of Computing Systems* 8.2 (2024). DOI: 10.1145/3656015 (cit. on p. 24).
 - [43] Dennis Trautwein et al. “Design and evaluation of IPFS: A storage layer for the decentralized web”. In: *Proceedings of the ACM SIGCOMM 2022 Conference*. 2022, pp. 739–752. DOI: 10.1145/3544216.3544232 (cit. on p. 24).
 - [44] Carmela Troncoso et al. “Systematizing decentralization and privacy: Lessons from 15 years of research and deployments”. In: *Proceedings on Privacy Enhancing Technologies* 2017.4 (2017), pp. 307–329. DOI: 10.1515/popets-2017-0056 (cit. on p. 24).
 - [45] Tao Zhou et al. “Bipartite network projection and personal recommendation”. In: *Physical Review E* 76.4 (2007) (cit. on p. 8).
 - [46] Lorenzo Zino and Ming Cao. “Analysis, prediction, and control of epidemics: A survey from scalar to dynamic network models”. In: *IEEE Circuits and Systems Magazine* 21.4 (2021). DOI: 10.1109/mcas.2021.3118100 (cit. on p. 17).